

EXPERIENCE

- AI Research Engineer

January 2026 - Present

Metafore

Bengaluru

– Built an end-to-end fine-tuning platform spanning dataset preparation, training, and deployment.

– Engineered ontology-aware knowledge graph workflows with LLM agents to structure and reason over domain data.

– Built a UI-driven agent generator that scaffolds production agents directly from Business Requirements Documents.
- Software Engineering Consultant

August 2025 - Jan 2026

Assured Health Technologies

Bengaluru

– Created 17 Python data-migration scripts to update fields across production database records.

– Automated AWS production deployments with CloudFormation, decreasing deployment time by 67%.

– Cut Django API latency from 35.7s to 3.7s using indexes and optimized prefetch queries.

– Adopted the Bun runtime to reduce application compile and execution time by 53%.
- Software Engineering Intern under LitmusChaos

July 2024 - September 2024

Cloud Native Computing Foundation Intern

Remote

– Built a Litmus upgrade agent that streamlined version transitions and reduced manual intervention by 30%.

– Engineered sequential upgrade pipelines for control-plane components and database migrations.

– Automated MongoDB schema upgrades for every Litmus release through versioned manifests.

– Optimized the Auth server login API, cutting latency 98.84% from 1.73s to 15.87ms.
- Google Summer Of Code (Intern at Organic Maps)

April 2023 - July 2023

Software Developer Intern

Remote

– Built a TypeScript/NodeJS proxy serving 50K+ daily requests at 99.9% uptime on Cloudflare Workers.

– Converted Google Maps URIs into mobile-ready custom formats, reducing application load time by 60%.

– Integrated the proxy across iOS and Android clients, improving mobile data-fetch speed by 30%.

PERSONAL PROJECTS

- LLM Inference Engine

A from-scratch LLM inference engine in PyTorch (KV cache, batching, paged memory)

Github

– Built a PyTorch LLM engine for GPT-2/Mistral with KV caching, continuous batching, and paged memory.

– Increased decode throughput 2.22x and cut CUDA time 46.6% with preallocated KV cache on A100.

– Scaled GPT-2 124M from 155 to 18,346 tokens/s (118.3x) at batch size 512 on A100.

– Cut peak KV memory 35.9%, reached 3.2x throughput at 64 concurrent requests, and added 122 tests.
- Kimi-K3-Impl

Reproduction of a 49M-parameter Kimi K3-inspired language model, trained on a single 8GB GPU

Github

– Reproduced a 49M-parameter Kimi K3 model with KDA, Gated MLA, Attention Residuals, and Stable LatentMoE.

– Trained on TinyStories in BF16 on one RTX 4060; selected step 9,750 at 1.6277 validation loss.

– Integrated Triton KDA kernels, cached KDA/MLA decoding, Muon optimizers, and FP32 CPU fallback.

– Added atomic checkpoints, live metrics, and 6 architecture, causality, decoding, and artifact tests.

TECHNICAL SKILLS AND INTERESTS

Languages: JavaScript, TypeScript, Go, Python, Rust, Java, C++, PHP, SQL
Frameworks: ExpressJS, ReactJS, Socket.io, Web3JS, jQuery
Cloud/Databases: AWS, Ansible, Docker, Prometheus, Apache Kafka, Thanos, OpenTelemetry, Kubernetes, MongoDB, MySQL
Coursework: Udemy certifications of Web Development and Kubernetes, completed Google Cloud Engineer track for Google Cloud event, GDSC IIT Patna, MA101 and MA102 (Linear algebra and real analysis)

EDUCATION

- Indian Institute of Technology Patna

2025

Bachelor of Technology