

- **Kimi-K3-Impl**

Reproduction of a 49M-parameter Kimi K3-inspired language model, trained on a single 8GB GPU

[Github](#)

- Reproduced a 49M-parameter Kimi K3 model with KDA, Gated MLA, Attention Residuals, and Stable LatentMoE layers.
- Trained on TinyStories in BF16 on one RTX 4060 Laptop GPU; selected final checkpoint 9,750 at 1.6277 validation loss.
- Integrated Triton KDA kernels, cached KDA/MLA autoregressive decoding, Muon optimizers, and an FP32 CPU fallback.
- Added atomic checkpoints, live training metrics, and 6 tests covering architecture, causality, decoding, and artifact integrity.