

- **Kimi-K3-Impl**

Reproduction of a 49M-parameter Kimi K3-inspired language model, trained on a single 8GB GPU

[Github](#)

- Reproduced a 49M-parameter Kimi K3 model with KDA, Gated MLA, Attention Residuals, and Stable LatentMoE.
- Trained on TinyStories in BF16 on one RTX 4060; selected step 9,750 at 1.6277 validation loss.
- Integrated Triton KDA kernels, cached KDA/MLA decoding, Muon optimizers, and FP32 CPU fallback.
- Added atomic checkpoints, live metrics, and 6 architecture, causality, decoding, and artifact tests.