

EXPERIENCE

- AI Research Engineer

January 2026 - Present

Metafore

Bengaluru

– Built an end-to-end fine-tuning platform spanning dataset preparation, model training, deployment, and downstream delivery.

– Engineered ontology-aware knowledge graph workflows with LLM agents to structure and reason over complex domain data.

– Built a UI-driven agent generator that scaffolds production agents directly from detailed Business Requirements Documents.

– Built MetaFabric’s Salesforce migration flow for metadata discovery, ontology mapping, extraction, validation, and lineage.
- Software Engineering Consultant

August 2025 - Jan 2026

Assured Health Technologies

Bengaluru

– Created and executed 17 Python data-migration scripts across the production database to update multiple application fields.

– Automated AWS deployments and operations with CloudFormation, decreasing overall end-to-end deployment time by 67%.

– Cut Django API latency from 35.7s to 3.7s using database indexes and eliminating redundant mutational prefetch queries.

– Adopted the Bun JavaScript runtime across application workflows, reducing combined compilation and runtime by 53%.
- Software Engineering Intern under LitmusChaos

July 2024 - September 2024

Cloud Native Computing Foundation Intern

Remote

– Built a Litmus upgrade agent that streamlined version transitions and reduced manual intervention by 30% during upgrades.

– Engineered sequential upgrade pipelines that reduced version-migration time for control-plane and Litmus database updates.

– Automated MongoDB schema migrations for every new Litmus version through release-specific deployment manifests.

– Optimized the Auth server login API by 98.84%, reducing measured login latency from 1.73 seconds to 15.87 milliseconds.
- Google Summer Of Code (Intern at Organic Maps)

April 2023 - July 2023

Software Developer Intern

Remote

– Built a TypeScript/NodeJS proxy serving 50K+ daily requests at 99.9% uptime on Cloudflare Workers edge infrastructure.

– Converted Google Maps URIs into mobile-ready custom formats, reducing application load time by 60% for mobile clients.

– Integrated the proxy across iOS and Android clients via streamlined API calls, improving mobile data-fetch speed by 30%.

PERSONAL PROJECTS

- LLM Inference Engine

A from-scratch LLM inference engine in PyTorch (KV cache, batching, paged memory)

Github

– Built a PyTorch LLM engine for GPT-2/Mistral with KV caching, continuous batching, paged memory, and async APIs.

– Increased decode throughput 2.22x and cut CUDA execution time 46.6% using a preallocated KV cache on one A100 GPU.

– Scaled GPT-2 124M throughput from 155 to 18,346 tokens/second, a 118.3x gain at batch size 512 on one A100 GPU.

– Cut KV-cache memory 35.9%, reached 3.2x throughput at 64 concurrent requests, and validated behavior with 122 tests.
- Kimi-K3-Impl

Reproduction of a 49M-parameter Kimi K3-inspired language model, trained on a single 8GB GPU

Github

– Reproduced a 49M-parameter Kimi K3 model with KDA, Gated MLA, Attention Residuals, and Stable LatentMoE layers.

– Trained on TinyStories in BF16 on one RTX 4060 Laptop GPU; selected final checkpoint 9,750 at 1.6277 validation loss.

– Integrated Triton KDA kernels, cached KDA/MLA autoregressive decoding, Muon optimizers, and an FP32 CPU fallback.

– Added atomic checkpoints, live training metrics, and 6 tests covering architecture, causality, decoding, and artifact integrity.

TECHNICAL SKILLS AND INTERESTS

Languages: JavaScript, TypeScript, Go, Python, Rust, Java, C++, PHP, SQL

Frameworks: ExpressJS, ReactJS, Socket.io, Web3JS, jQuery

Cloud/Databases: AWS, Ansible, Docker, Prometheus, Apache Kafka, Thanos, OpenTelemetry, Kubernetes, MongoDB, MySQL

Coursework: Udemy certifications of Web Development and Kubernetes, completed Google Cloud Engineer track for Google Cloud event, GDSC IIT Patna, MA101 and MA102 (Linear algebra and real analysis)

EDUCATION

- Indian Institute of Technology Patna

2025

Bachelor of Technology