

EXPERIENCE

- AI Research Engineer

January 2026 - Present

Metafore

Bengaluru

– Built an end-to-end model fine-tuning platform, streamlining the pipeline from dataset preparation through training to deployment for downstream teams.

– Engineered ontology-aware knowledge graph workflows integrated with LLM agents to structure and reason over domain data.

– Developed a UI-driven agent generator that scaffolds production agents directly from a Business Requirements Document (BRD).
- Software Engineering Consultant

August 2025 - Jan 2026

Assured Health Technologies

Bengaluru

– Created 17 data migration scripts in Python that ran across the database to update various fields.

– Used AWS for deployment and automated production operations using CloudFormation, decreasing deployment time by 67%.

– Optimized Django API to reduce time from 35.7s to 3.7s using indexes and removing extra mutational queries on prefetching.

– Utilized Bun runtime to reduce compile and run time by 53%.
- Software Engineering Intern under LitmusChaos

July 2024 - September 2024

Cloud Native Computing Foundation Intern

Remote

– Developed an upgrade agent streamlining version transitions for Litmus, leading to a 30% reduction in manual intervention time during upgrades and enhancing overall user experience across the platform.

– Engineered sequential upgrade pipeline reducing version migration time for control plane and database updates.

– Automated the process of upgrading the MongoDB schema with every new version of Litmus via manifests.

– Optimized the login API of the Auth server by 98.84% by bringing down the latency from 1.73s to 15.87ms.
- Google Summer Of Code (Intern at Organic Maps)

April 2023 - July 2023

Software Developer Intern

Remote

– Built high-performance proxy using TypeScript and NodeJS, serving 50K+ requests/day with 99.9% uptime on Cloudflare Workers.

– Architected proxy to convert Google Maps URIs to custom formats, reducing load time by 60%.

– Integrated the newly prepared proxy seamlessly into both iOS and Android applications, enabling a streamlined API interaction that improved data fetching speed by 30% across mobile platforms.

PERSONAL PROJECTS

- LLM Inference Engine

A from-scratch LLM inference engine in PyTorch (KV cache, batching, paged memory)

Github

– Built a from-scratch LLM inference engine in PyTorch with KV caching, continuous batching, and paged memory.

– Achieved 2.22× decode speedup from KV caching at 1K tokens.

– Gained 118× throughput via batched inference on A100; cut memory 8.7× with a paged cache.

– Hit 11.1× serving throughput under 64 concurrent users; validated across 122 tests.
- gogent

A Go framework for LLM agents with built-in OpenTelemetry observability

Github

– Designed a Go framework for building LLM agents, providing a structured runtime for tool use, provider integration, and agent orchestration.

– Built in OpenTelemetry observability to emit traces and metrics for agent runs, making multi-step agent pipelines debuggable and monitorable.

– Abstracted LLM provider access behind a common interface, supporting multiple providers and enabling drop-in swapping without changing agent code.

– Included a memory subsystem and examples suite to demonstrate reusable agent patterns out of the box.

TECHNICAL SKILLS AND INTERESTS

Languages: JavaScript, TypeScript, Go, Python, Rust, C++, PHP, SQL
Frameworks: ExpressJS, ReactJS, Socket.io, Web3JS, jQuery
Cloud/Databases: AWS, Ansible, Docker, Prometheus, Apache Kafka, Thanos, OpenTelemetry, Kubernetes, MongoDB, MySQL
Coursework: Udemy certifications of Web Development and Kubernetes, completed Google Cloud Engineer track for Google Cloud event, GDSC IIT Patna, MA101 and MA102 (Linear algebra and real analysis)

EDUCATION

- Indian Institute of Technology Patna

2025

Bachelor of Technology